

Image Colorization with CNN

```
In [ ]: import numpy as np
import pandas as pd # used for data processing and reading csv
import os
import cv2
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from tensorflow.keras.optimizers import Adam
from keras import backend as K
from keras.layers import Conv2D,MaxPooling2D,UpSampling2D,Input,BatchN
from keras.layers.merge import concatenate
from keras.models import Model
from keras.preprocessing.image import ImageDataGenerator
import keras
#from tensorflow import set_random_seed
import tensorflow.compat.v1 as tf

tf.set_random_seed(123)
session_conf = tf.ConfigProto(intra_op_parallelism_threads=1, inter_op
sess = tf.Session(graph=tf.get_default_graph(), config=session_conf)
tf.keras.backend.set_session(sess)
tf.set_random_seed(2)
np.random.seed(1)
from google.colab import drive
drive.mount('/content/drive')
```

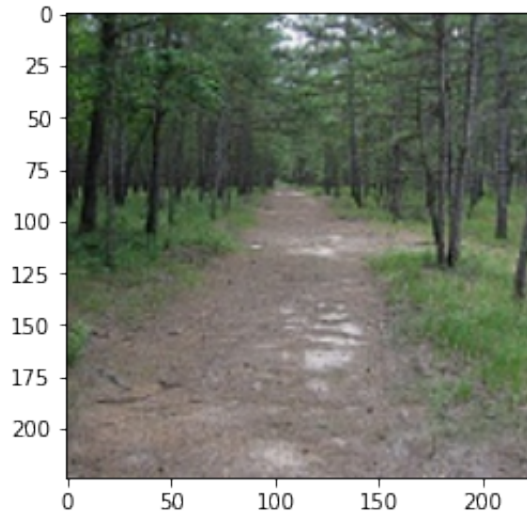
Mounted at /content/drive

Loading the image path

```
In [ ]: ImagePath="/content/drive/MyDrive/Deep Learning/"
```

```
In [ ]: img = cv2.imread(ImagePath+"20056.jpg")
img = cv2.cvtColor(img,cv2.COLOR_BGR2RGB)
img = cv2.resize(img, (224, 224))
plt.imshow(img)
img.shape
```

Out[4]: (224, 224, 3)



```
In [ ]: HEIGHT=224
        WIDTH=224
        ImagePath="/content/drive/MyDrive/Deep Learning/"

        def ExtractInput(path):
            X_img=[]
            y_img=[]
            for imageDir in os.listdir(ImagePath):
                try:
                    img = cv2.imread(ImagePath + imageDir)
                    img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
                    img = cv2.cvtColor(img, cv2.COLOR_RGB2GRAY)
                    img = cv2.cvtColor(img, cv2.COLOR_GRAY2RGB)

                    img = img.astype(np.float32)
                    img_lab = cv2.cvtColor(img, cv2.COLOR_RGB2Lab)
                    img_lab_rs = cv2.resize(img_lab, (WIDTH, HEIGHT)) # resize
                    img_l = img_lab_rs[:, :, 0] # pull out L channel
                    img_ab = img_lab_rs[:, :, 1:] # Extracting the ab channel
                    img_ab = img_ab/128
                    X_img.append(img_l)
                    y_img.append(img_ab)
                except:
                    pass
            X_img = np.array(X_img)
            y_img = np.array(y_img)

            return X_img, y_img
```

```
In [ ]: X_, y_ = ExtractInput(ImagePath) # Data-preprocessing
```

```
In [ ]: X_train, X_val, y_train, y_val = train_test_split(X_, y_, random_state=42)
```

```
In [ ]: print()
```

```

In [ ]: K.clear_session()
def InstantiateModel(in_):
    model_ = Conv2D(16,(3,3),padding='same',strides=1)(in_)
    model_ = LeakyReLU()(model_)
    model_ = Conv2D(32,(3,3),padding='same',strides=1)(model_)
    model_ = LeakyReLU()(model_)
    model_ = BatchNormalization()(model_)
    model_ = MaxPooling2D(pool_size=(2,2),padding='same')(model_)

    model_ = Conv2D(64,(3,3),padding='same',strides=1)(model_)
    model_ = LeakyReLU()(model_)
    model_ = BatchNormalization()(model_)
    model_ = MaxPooling2D(pool_size=(2,2),padding='same')(model_)

    model_ = Conv2D(128,(3,3),padding='same',strides=1)(model_)
    model_ = LeakyReLU()(model_)
    model_ = BatchNormalization()(model_)

    model_ = Conv2D(256,(3,3),padding='same',strides=1)(model_)
    model_ = LeakyReLU()(model_)
    model_ = BatchNormalization()(model_)

    model_ = UpSampling2D((2, 2))(model_)
    model_ = Conv2D(128,(3,3),padding='same',strides=1)(model_)
    model_ = LeakyReLU()(model_)
    model_ = BatchNormalization()(model_)

    model_ = UpSampling2D((2, 2))(model_)
    model_ = Conv2D(64,(3,3),padding='same',strides=1)(model_)
    model_ = LeakyReLU()(model_)

    concat_ = concatenate([model_, in_])

    model_ = Conv2D(64,(3,3),padding='same',strides=1)(concat_)
    model_ = LeakyReLU()(model_)
    model_ = BatchNormalization()(model_)

    model_ = Conv2D(32,(3,3),padding='same',strides=1)(model_)
    model_ = LeakyReLU()(model_)

    model_ = Conv2D(2,(3,3),activation='tanh',padding='same',strides=1)(model_)

    return model_

```

In []:

```

Input_Sample = Input(shape=(HEIGHT, WIDTH,1))
Output_ = InstantiateModel(Input_Sample)
Model_Colourization = Model(inputs=Input_Sample, outputs=Output_)

```

In []:

```

LEARNING_RATE = 0.001
Model_Colourization.compile(optimizer=Adam(lr=LEARNING_RATE),
                             loss='mean_squared_error')
Model_Colourization.summary()

```

Model: "model"

Layer (type) ected to	Output Shape	Param #	Conn
=====			
input_1 (InputLayer)	[(None, 224, 224, 1 0)]		[]
conv2d (Conv2D) put_1[0][0]']	(None, 224, 224, 16 160)		['in
leaky_re_lu (LeakyReLU) nv2d[0][0]']	(None, 224, 224, 16 0)		['co
conv2d_1 (Conv2D) aky_re_lu[0][0]']	(None, 224, 224, 32 4640)		['le
leaky_re_lu_1 (LeakyReLU) nv2d_1[0][0]']	(None, 224, 224, 32 0)		['co
batch_normalization (BatchNorm aky_re_lu_1[0][0]'] alization)	(None, 224, 224, 32 128)		['le
max_pooling2d (MaxPooling2D) tch_normalization[0][0]']	(None, 112, 112, 32 0)		['ba

```

conv2d_2 (Conv2D) (None, 112, 112, 64 18496 ['ma
x_pooling2d[0][0]']
)

leaky_re_lu_2 (LeakyReLU) (None, 112, 112, 64 0 ['co
nv2d_2[0][0]']
)

batch_normalization_1 (BatchNo (None, 112, 112, 64 256 ['le
aky_re_lu_2[0][0]']
rmalization)
)

max_pooling2d_1 (MaxPooling2D) (None, 56, 56, 64) 0 ['ba
tch_normalization_1[0][0]']

conv2d_3 (Conv2D) (None, 56, 56, 128) 73856 ['ma
x_pooling2d_1[0][0]']

leaky_re_lu_3 (LeakyReLU) (None, 56, 56, 128) 0 ['co
nv2d_3[0][0]']

batch_normalization_2 (BatchNo (None, 56, 56, 128) 512 ['le
aky_re_lu_3[0][0]']
rmalization)

conv2d_4 (Conv2D) (None, 56, 56, 256) 295168 ['ba
tch_normalization_2[0][0]']

leaky_re_lu_4 (LeakyReLU) (None, 56, 56, 256) 0 ['co
nv2d_4[0][0]']

batch_normalization_3 (BatchNo (None, 56, 56, 256) 1024 ['le
aky_re_lu_4[0][0]']
rmalization)

up_sampling2d (UpSampling2D) (None, 112, 112, 25 0 ['ba
tch_normalization_3[0][0]']
6)

conv2d_5 (Conv2D) (None, 112, 112, 12 295040 ['up
_sampling2d[0][0]']
8)

leaky_re_lu_5 (LeakyReLU) (None, 112, 112, 12 0 ['co
nv2d_5[0][0]']
8)

batch_normalization_4 (BatchNo (None, 112, 112, 12 512 ['le
aky_re_lu_5[0][0]']
rmalization)
8)

```

```

up_sampling2d_1 (UpSampling2D) (None, 224, 224, 12 0 ['ba
tch_normalization_4[0][0]']
8)

conv2d_6 (Conv2D) (None, 224, 224, 64 73792 ['up
_sampling2d_1[0][0]']
)

leaky_re_lu_6 (LeakyReLU) (None, 224, 224, 64 0 ['co
nv2d_6[0][0]']
)

concatenate (Concatenate) (None, 224, 224, 65 0 ['le
aky_re_lu_6[0][0]',
)
'input_1[0][0]']

conv2d_7 (Conv2D) (None, 224, 224, 64 37504 ['co
ncatenate[0][0]']
)

leaky_re_lu_7 (LeakyReLU) (None, 224, 224, 64 0 ['co
nv2d_7[0][0]']
)

batch_normalization_5 (BatchNo (None, 224, 224, 64 256 ['le
aky_re_lu_7[0][0]']
rmalization)
)

conv2d_8 (Conv2D) (None, 224, 224, 32 18464 ['ba
tch_normalization_5[0][0]']
)

leaky_re_lu_8 (LeakyReLU) (None, 224, 224, 32 0 ['co
nv2d_8[0][0]']
)

conv2d_9 (Conv2D) (None, 224, 224, 2) 578 ['le
aky_re_lu_8[0][0]']

```

```

=====
=====

```

```

Total params: 820,386
Trainable params: 819,042
Non-trainable params: 1,344

```

```

/usr/local/lib/python3.7/dist-packages/keras/optimizer_v2/adam.py:105
: UserWarning: The `lr` argument is deprecated. Use `learning rate` i

```

• Observing: The `learning_rate` argument is deprecated, use `learning_rate_decay` instead.

```
super(Adam, self).__init__(name, **kwargs)
```

```
In [ ]: def GenerateInputs(X_,y_):
        for i in range(len(X_)):
            X_input = X_[i].reshape(1,224,224,1)
            y_input = y_[i].reshape(1,224,224,2)
            yield (X_input,y_input)
checkpoint_filepath = '/content/drive/MyDrive/CNN/checkpoint'
model_checkpoint_callback = tf.keras.callbacks.ModelCheckpoint(
    filepath=checkpoint_filepath,
    save_weights_only=True,
    monitor='val_acc',
    mode='max',
    save_best_only=True)
Model_Colourization = keras.models.load_model('/content/drive/MyDrive/
Model_Colourization.fit_generator(GenerateInputs(X_,y_),epochs=40,verb
Model_Colourization.save('/content/drive/MyDrive/CNN/Model_Colourizati
```

Epoch 1/40

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:14: User Warning: `Model.fit_generator` is deprecated and will be removed in a future version. Please use `Model.fit`, which supports generators.

40/40 [=====] - 80s 2s/step - loss: 6.2425e-06

Epoch 2/40

40/40 [=====] - 74s 2s/step - loss: 5.5684e-06

Epoch 3/40

40/40 [=====] - 73s 2s/step - loss: 1.9193e-06

Epoch 4/40

40/40 [=====] - 73s 2s/step - loss: 1.4038e-05

Epoch 5/40

40/40 [=====] - 73s 2s/step - loss: 6.8002e-06

Epoch 6/40

40/40 [=====] - 73s 2s/step - loss: 1.5060e-06

06

Epoch 7/40

40/40 [=====] - 73s 2s/step - loss: 9.5100e-06

Epoch 8/40

40/40 [=====] - 72s 2s/step - loss: 3.9833e-06

Epoch 9/40


```
Epoch 9/40
40/40 [=====] - 73s 2s/step - loss: 9.4687e-06
Epoch 10/40
40/40 [=====] - 73s 2s/step - loss: 1.1352e-05
Epoch 11/40
40/40 [=====] - 73s 2s/step - loss: 6.8041e-06
Epoch 12/40
40/40 [=====] - 73s 2s/step - loss: 1.8981e-05
Epoch 13/40
40/40 [=====] - 73s 2s/step - loss: 1.7091e-05
Epoch 14/40
40/40 [=====] - 73s 2s/step - loss: 6.8975e-06
Epoch 15/40
40/40 [=====] - 73s 2s/step - loss: 2.9812e-06
Epoch 16/40
40/40 [=====] - 73s 2s/step - loss: 2.4602e-06
Epoch 17/40
40/40 [=====] - 73s 2s/step - loss: 3.0610e-06
Epoch 18/40
40/40 [=====] - 73s 2s/step - loss: 8.2604e-06
Epoch 19/40
40/40 [=====] - 73s 2s/step - loss: 2.7190e-05
Epoch 20/40
40/40 [=====] - 73s 2s/step - loss: 3.3542e-06
Epoch 21/40
40/40 [=====] - 73s 2s/step - loss: 4.9227e-06
Epoch 22/40
40/40 [=====] - 73s 2s/step - loss: 1.1562e-05
Epoch 23/40
40/40 [=====] - 73s 2s/step - loss: 1.0428e-05
Epoch 24/40
40/40 [=====] - 72s 2s/step - loss: 7.1335e-06
Epoch 25/40
40/40 [=====] - 73s 2s/step - loss: 7.2008e-06
```

uu

Epoch 26/40

40/40 [=====] - 73s 2s/step - loss: 1.1901e-05

Epoch 27/40

40/40 [=====] - 73s 2s/step - loss: 1.1012e-05

Epoch 28/40

40/40 [=====] - 73s 2s/step - loss: 2.0832e-06

Epoch 29/40

40/40 [=====] - 73s 2s/step - loss: 2.4532e-06

Epoch 30/40

40/40 [=====] - 73s 2s/step - loss: 1.1199e-05

Epoch 31/40

40/40 [=====] - 73s 2s/step - loss: 7.7890e-05

Epoch 32/40

40/40 [=====] - 73s 2s/step - loss: 4.7251e-05

Epoch 33/40

40/40 [=====] - 73s 2s/step - loss: 2.3196e-06

Epoch 34/40

40/40 [=====] - 73s 2s/step - loss: 1.2447e-06

Epoch 35/40

40/40 [=====] - 73s 2s/step - loss: 2.0690e-06

Epoch 36/40

40/40 [=====] - 73s 2s/step - loss: 8.4212e-07

Epoch 37/40

40/40 [=====] - 73s 2s/step - loss: 6.4166e-07

Epoch 38/40

40/40 [=====] - 73s 2s/step - loss: 5.5359e-07

Epoch 39/40

40/40 [=====] - 73s 2s/step - loss: 3.4691e-06

Epoch 40/40

40/40 [=====] - 73s 2s/step - loss: 1.5790e-06

In []:

```
TestImagePath="/content/drive/MyDrive/Deep Learning/"
```

In []:

```
def ExtractTestInput(ImagePath):  
    img = cv2.imread(ImagePath)  
    img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)  
    img = cv2.cvtColor(img, cv2.COLOR_RGB2GRAY)  
    img_ = cv2.cvtColor(img, cv2.COLOR_GRAY2RGB)  
    img_ = cv2.cvtColor(img_, cv2.COLOR_RGB2Lab)  
    img_ = img_.astype(np.float32)  
    img_lab_rs = cv2.resize(img_, (WIDTH, HEIGHT)) # resize image to r  
    img_l = img_lab_rs[:, :, 0] # pull out L channel  
    img_l_reshaped = img_l.reshape(1, 224, 224, 1)  
  
    return img_l_reshaped
```

In []:

```
ImagePath=TestImagePath+"20074.jpg"  
image_for_test = ExtractTestInput(ImagePath)  
Prediction = Model_Colourization.predict(image_for_test)  
Prediction = Prediction*128  
Prediction=Prediction.reshape(224, 224, 2)
```

```

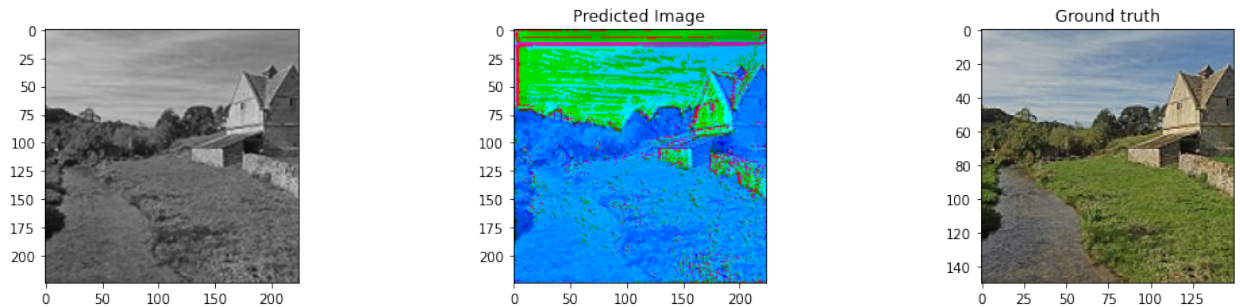
In [ ]: plt.figure(figsize=(30,20))
plt.subplot(5,5,1)
img = cv2.imread(TestImagePath+"20074.jpg")
img_1 = cv2.cvtColor(img,cv2.COLOR_BGR2RGB)
img = cv2.cvtColor(img_1, cv2.COLOR_RGB2GRAY)
img = cv2.cvtColor(img, cv2.COLOR_GRAY2RGB)
img = cv2.resize(img, (224, 224))
plt.imshow(img)

plt.subplot(5,5,1+1)
img_ = cv2.cvtColor(img, cv2.COLOR_RGB2Lab)
img_[:,:,1:] = Prediction
img_ = cv2.cvtColor(img_, cv2.COLOR_Lab2RGB)
plt.title("Predicted Image")
plt.imshow(img_)

plt.subplot(5,5,1+2)
plt.title("Ground truth")
plt.imshow(img_1)

```

Out[21]: <matplotlib.image.AxesImage at 0x7fd8c48b7150>



```

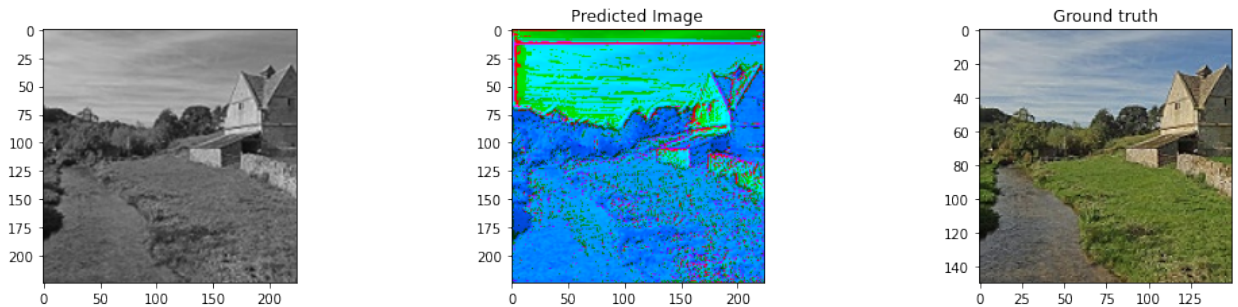
In [ ]: plt.figure(figsize=(30,20))
plt.subplot(5,5,1)
img = cv2.imread(TestImagePath+"20074.jpg")
img_1 = cv2.cvtColor(img,cv2.COLOR_BGR2RGB)
img = cv2.cvtColor(img_1, cv2.COLOR_RGB2GRAY)
img = cv2.cvtColor(img, cv2.COLOR_GRAY2RGB)
img = cv2.resize(img, (224, 224))
plt.imshow(img)

plt.subplot(5,5,1+1)
img_ = cv2.cvtColor(img, cv2.COLOR_RGB2Lab)
img_[:,:,1:] = Prediction
img_ = cv2.cvtColor(img_, cv2.COLOR_Lab2RGB)
plt.title("Predicted Image")
plt.imshow(img_)

plt.subplot(5,5,1+2)
plt.title("Ground truth")
plt.imshow(img_1)

```

Out[16]: <matplotlib.image.AxesImage at 0x7fd8c48f6dd0>



```

In [ ]: ImagePath=TestImagePath+"20080.jpg"
image_for_test = ExtractTestInput(ImagePath)
Prediction_1 = Model_Colourization.predict(image_for_test)
Prediction_1 = Prediction_1*128
Prediction_1=Prediction_1.reshape(224,224,2)

```

```


In [ ]: plt.figure(figsize=(30,20))
plt.subplot(5,5,1)
img = cv2.imread(TestImagePath+"20080.jpg")
img_1 = cv2.cvtColor(img,cv2.COLOR_BGR2RGB)
img = cv2.cvtColor(img_1, cv2.COLOR_RGB2GRAY)
img = cv2.cvtColor(img, cv2.COLOR_GRAY2RGB)
img = cv2.resize(img, (224, 224))
plt.imshow(img)

plt.subplot(5,5,1+1)
img_ = cv2.cvtColor(img, cv2.COLOR_RGB2Lab)
img_[:,:,1:] = Prediction_1
img_ = cv2.cvtColor(img_, cv2.COLOR_Lab2RGB)
plt.title("Predicted Image")
plt.imshow(img_)

plt.subplot(5,5,1+2)
plt.title("Ground truth")
plt.imshow(img_1)

```

Out[56]: <matplotlib.image.AxesImage at 0x7fe76d630250>



```

In [ ]: ImagePath=TestImagePath+"20147.jpg"
image_for_test = ExtractTestInput(ImagePath)
Prediction_2 = Model_Colourization.predict(image_for_test)
Prediction_2 = Prediction_2*128
Prediction_2=Prediction_2.reshape(224,224,2)

```

```

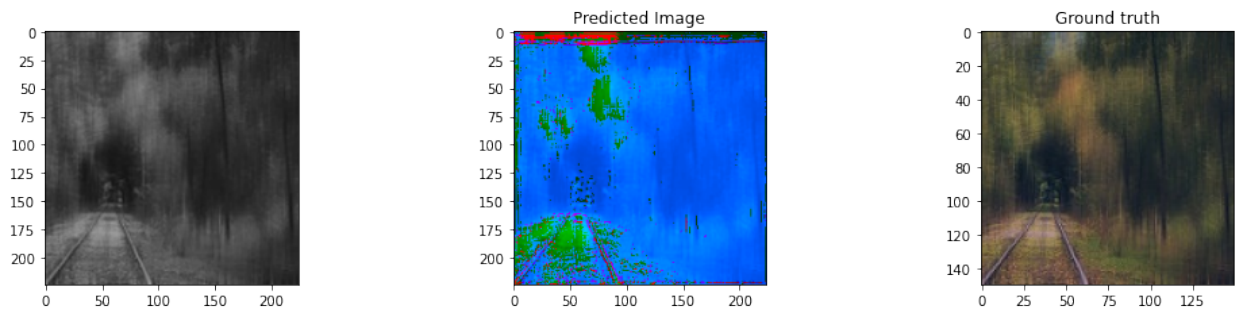
In [ ]: plt.figure(figsize=(30,20))
plt.subplot(5,5,1)
img = cv2.imread(TestImagePath+"20147.jpg")
img_1 = cv2.cvtColor(img,cv2.COLOR_BGR2RGB)
img = cv2.cvtColor(img_1, cv2.COLOR_RGB2GRAY)
img = cv2.cvtColor(img, cv2.COLOR_GRAY2RGB)
img = cv2.resize(img, (224, 224))
plt.imshow(img)

plt.subplot(5,5,1+1)
img_ = cv2.cvtColor(img, cv2.COLOR_RGB2Lab)
img_[:,:,1:] = Prediction_2
img_ = cv2.cvtColor(img_, cv2.COLOR_Lab2RGB)
plt.title("Predicted Image")
plt.imshow(img_)

plt.subplot(5,5,1+2)
plt.title("Ground truth")
plt.imshow(img_1)

```

Out[58]: <matplotlib.image.AxesImage at 0x7fe76d4e9a50>



```

In [ ]: TestImagePath="/content/drive/MyDrive/Deep Learning/21461.jpg"
image_for_test = ExtractTestInput(TestImagePath)
Prediction_3 = Model_Colourization.predict(image_for_test)
Prediction_3 = Prediction_3*128
Prediction_3=Prediction_3.reshape(224,224,2)

```

```


In [ ]: plt.figure(figsize=(30,20))
plt.subplot(5,5,1)
img = cv2.imread(TestImagePath)
img_1 = cv2.cvtColor(img,cv2.COLOR_BGR2RGB)
img = cv2.cvtColor(img_1, cv2.COLOR_RGB2GRAY)
img = cv2.cvtColor(img, cv2.COLOR_GRAY2RGB)
img = cv2.resize(img, (224, 224))
plt.imshow(img)

plt.subplot(5,5,1+1)
img_ = cv2.cvtColor(img, cv2.COLOR_RGB2Lab)
img[:, :, 1:] = Prediction_3
img_ = cv2.cvtColor(img_, cv2.COLOR_Lab2RGB)
plt.title("Predicted Image")
plt.imshow(img_)

plt.subplot(5,5,1+2)
plt.title("Ground truth")
plt.imshow(img_1)

```

Out[60]: <matplotlib.image.AxesImage at 0x7fe76d3b1210>



```

In [ ]: TestImagePath="/content/drive/MyDrive/Deep Learning/21422.jpg"
image_for_test = ExtractTestInput(TestImagePath)
Prediction_4 = Model_Colourization.predict(image_for_test)
Prediction_4 = Prediction_4*128
Prediction_4=Prediction_4.reshape(224,224,2)

```



```


In [ ]: plt.figure(figsize=(30,20))
plt.subplot(5,5,1)
img = cv2.imread(TestImagePath)
img_1 = cv2.cvtColor(img,cv2.COLOR_BGR2RGB)
img = cv2.cvtColor(img_1, cv2.COLOR_RGB2GRAY)
img = cv2.cvtColor(img, cv2.COLOR_GRAY2RGB)
img = cv2.resize(img, (224, 224))
plt.imshow(img)

plt.subplot(5,5,1+1)
img_ = cv2.cvtColor(img, cv2.COLOR_RGB2Lab)
img_[:,:,1:] = Prediction_4
img_ = cv2.cvtColor(img_, cv2.COLOR_Lab2RGB)
plt.title("Predicted Image")
plt.imshow(img_)

plt.subplot(5,5,1+2)
plt.title("Ground truth")
plt.imshow(img_1)

```

Out[62]: <matplotlib.image.AxesImage at 0x7fe76c3ad9d0>



```

In [ ]: TestImagePath="/content/drive/MyDrive/Deep Learning/21486.jpg"
image_for_test = ExtractTestInput(TestImagePath)
Prediction_5 = Model_Colourization.predict(image_for_test)
Prediction_5 = Prediction_5*128
Prediction_5=Prediction_5.reshape(224,224,2)

```

```

In [ ]: plt.figure(figsize=(30,20))
plt.subplot(5,5,1)
img = cv2.imread(TestImagePath)
img_1 = cv2.cvtColor(img,cv2.COLOR_BGR2RGB)
img = cv2.cvtColor(img_1, cv2.COLOR_RGB2GRAY)
img = cv2.cvtColor(img, cv2.COLOR_GRAY2RGB)
img = cv2.resize(img, (224, 224))
plt.imshow(img)

plt.subplot(5,5,1+1)
img_ = cv2.cvtColor(img, cv2.COLOR_RGB2Lab)
img_[:,:,1:] = Prediction_5
img_ = cv2.cvtColor(img_, cv2.COLOR_Lab2RGB)
plt.title("Predicted Image")
plt.imshow(img_)

plt.subplot(5,5,1+2)
plt.title("Ground truth")
plt.imshow(img_1)

```

Out[64]: <matplotlib.image.AxesImage at 0x7fe76b68de90>

